

MVP: the same attribution on Blocksworld

2026

This is the same column-wise procedure on a second search. Four blocks, beam 4, depth 4, forty episodes, drift at episode 16, judge mix 0.30. $Y = 1$ when the goal tower (A, B, C, D) is returned. The state encoding is four numbers in $[0, 1]$: stack count, tallest stack, goal-prefix length, and whether D is clear. No text encoder, no preset layer weights, no AUC.

Y moves from 1.00 to 0.542.

Layer	SSE share	Best column	Best p	Largest mean shift
State encoding	0.560	stack count e_0	0.012	0.115 on tallest stack
Graph	0.293	children	0.012	0.578 on children
Statistics	0.060	score z	0.025	1.051 on score z
Interaction	0.059	children per depth	0.025	0.242
History	0.018	rolling std	0.160	0.354 on previous Y
Context	0.009	last-step fraction	0.247	0.005

On Game of 24 the graph layer was 9% of the SSE and the score profile was the larger part. Here the committed tower encoding is 56% and the child count is 29%, and both clear $p = 0.012$. The drifted judge is expanding a different frontier: more or fewer legal moves, and a different tower shape. History does not clear 0.10. Beam and maximum depth stay fixed, so context stays near zero.

The screen is unchanged. One column, one SSE drop, one mean shift. Blocksworld only changes which columns are allowed to be nonzero.